

The Payne

For this particular case, we generate ~160 training spectra and ~160 testing spectra using B. Kurucz codes and P. Cargile line list. The Teff and log g are drawn from the MIST isochrones. We consider stellar evolution from the main sequence stars up to the red clump stars with Teff from 3500 K to 7000 K. We also consider a wide range of metallicity with $[Fe/H] = -2$ to 0.5.

The following will show that using only 160 models, The Payne is an extremely precise generative model for this wide range of stellar parameters.

Train neural network

This training step is quite slow. Training a single wavelength pixel typically takes 10 minutes to 1 hour. It requires running on a cluster. If you want to try this part on a personal laptop, make sure that you only train spectra with a few wavelength pixels. Note that although the training is excruciating slow, once the network is trained, fitting spectra (i.e. testing step) is very efficient.

```
In [ ]: # import python packages
import numpy as np
from multiprocessing import Pool

import theano
import theano.tensor as T
from theano.tensor.nnet import sigmoid

#=====
# number of CPUs for parallel computing
num_CPU = 4

#-----
### some parameters of neural network ###
# how many portions will we split the training data to perform stochastic gradient descent
# smaller batch size (i.e. larger number) will take longer to converge
```

```

# smaller batch size (i.e. larger number) will take longer to converge
# but might converge with a smaller number of steps
mini_batch_size = 2

# initial step size in stochastic gradient descent
# smaller step size will be slower but will provide better convergence
eta_choice = 0.1

# the minimum step size beyond which we will truncate
min_eta = 0.001

# how many steps of gradient descent per loop are we going to perform
num_epochs_choice = 10000

# truncation criterion
trunc_diff = 0.003

# maximum number of loop (to avoid infinite loop)
# i.e. max_iter*num_epochs_choice is the maximum number of steps beyond which we will truncate
max_iter = 100

# how many neurons per layer
# here we always consider two fully connected layers
n_neurons = 10

#-----
# define activation function that we will use in the validation step
# make sure this function is consistent with the training function
# here we choose a sigmoid function
def act_func(z):
    return 1.0/(1.0+np.exp(-z))

#=====
# restore all spectra
# the variable "spectra" has a dimension of (# pixels, # spectra)
# the variable "labels" has a dimension of (# labels, # spectra)
# for this particular training set, the labels are (teff[K], logg, Fe/H)

```

```

temp = np.load("kurucz_spectra.npz")
spectra = temp["Y_u_all"].T
labels = temp["labels_array"]

# only use half of them to train the neural network
# we will save the other half for testing
spectra = spectra[:, ::2]
labels = labels[:, ::2]

#-----
# neural networks typically train a function mapping from [0,1] -> [0,1]
# so here we scale both input (labels) and output (fluxes) to [0.1,0.9]

# record how we scale the labels
x_max = np.max(labels, axis=1)
x_min = np.min(labels, axis=1)

# scale the labels
labels = ((labels.T - x_min)*0.8/(x_max-x_min) + 0.1).T

# scale the fluxes
# here we assume normalized spectra
# for flux spectra this scaling has to be changed
spectra = spectra*0.8 + 0.1

#=====
# theano is a very powerful package to train neural network
# it performs "auto diff", i.e., provides analytic differentiation of any cost function

# convert labels into a theano variable
training_x = theano.shared(np.asarray(labels.T, dtype='float64'))

#-----
# main network class
class Network(object):
    def init (self, layers, mini batch size):

```

```

# initiate network properties
self.layers = layers
self.mini_batch_size = mini_batch_size
self.params = [param for layer in self.layers for param in layer.params]

self.x = T.dmatrix("x")
self.y = T.dmatrix("y")

init_layer = self.layers[0]
init_layer.set_inpt(self.x, self)
for j in xrange(1, len(self.layers)):
    prev_layer, layer = self.layers[j-1], self.layers[j]
    layer.set_inpt(prev_layer.output, self)
self.output = self.layers[-1].output

# create a property to record the cost function at each training step
self.cost_train = []

```

```

#-----
# stochastic gradient descent
def SGD(self, training_x, training_y, epochs, eta):

    # reset the cost for each training loop
    self.cost_train = []

    # compute mini batches
    num_sample = training_x.get_value().shape[0]
    num_training_batches = num_sample/self.mini_batch_size

    # define cost function, symbolic gradients, and updates
    cost = self.layers[-1].cost(self)
    grads = T.grad(cost, self.params)
    updates = [(param, param-eta/self.mini_batch_size*grad)\
                for param, grad in zip(self.params, grads)]

    i = T.lscalar()

```

```

# randomize the training data for stochastic gradient descent
ind = np.arange(num_sample)
np.random.shuffle(ind)
ind = theano.shared(ind)

# define function to train a mini-batch
train_mb = theano.function([i], cost, updates=updates,
    givens={self.x: training_x[ind[i*self.mini_batch_size:(i+1)*self.mini_batch_size]],
            self.y: training_y[ind[i*self.mini_batch_size:(i+1)*self.mini_batch_size]]})

# the actual training
for epoch in xrange(epochs):
    cost_train_ij = 0.
    for minibatch_index in xrange(num_training_batches):
        # sum up all cost for each mini batch
        cost_train_ij += train_mb(minibatch_index)
    self.cost_train.append(cost_train_ij)

#-----
# define fully connected layers
### here we choose sigmoid function to be the activation function ###
class FullyConnectedLayer(object):
    def __init__(self, n_in, n_out, activation_fn=sigmoid):
        self.n_in = n_in
        self.n_out = n_out
        self.activation_fn = activation_fn

# initialize weights and biases of the neural network
self.w = theano.shared(np.asarray(np.random.normal(
    loc=0.0, scale=np.sqrt(1.0/n_out), size=(n_in,n_out)),
    dtype=theano.config.floatX))
self.b = theano.shared(np.asarray(np.random.normal(\
    loc=0.0, scale=1.0, size=(n_out,)),
    dtype=theano.config.floatX))
self.params = [self.w,self.b]

```

```

#-----
# define input and output for each neural network layer
def set_inpt(self, inpt, net):
    self.inpt = inpt.reshape((net.mini_batch_size, self.n_in))
    self.output = self.activation_fn(T.dot(self.inpt, self.w) + self.b)

### define a cost function ###
def cost(self, net):
    return T.sum(T.abs_(net.y-self.output))

#=====
# define training function for each wavelength pixel to run in parallel
# note we create individual neural network for each wavelength pixel
def train_pixel(pixel_no):

    # extract flux of a wavelength pixel
    training_y = theano.shared(np.asarray(np.array([spectra[pixel_no,:]]).T, dtype='float64'))

    # define the network
    net = Network([
        FullyConnectedLayer(n_in=training_x.get_value().shape[1], n_out=n_neurons),\
        FullyConnectedLayer(n_in=n_neurons, n_out=training_y.get_value().shape[1]),\
        mini_batch_size)

#-----

# initiate loop counter and step size
loop_count = 0
step_devide = 1.

# sometimes the network can stuck at the initial point
# so first we train for 1000 steps
net.SGD(training_x, training_y, 1000, eta_choice)

# we evaluate if the cost has improved
while (np.abs(np.mean(net.cost_train[:100]) \
               - np.mean(net.cost_train[-100:]))/(np.mean(net.cost_train[:100])) < 0.1)\
    and (loop_count < max_iter):

```

```

        and (loop_count < max_iter):

# if not we reset the newtwork (and hence the initial point)
# and loop until it finds a valid initial point
net = Network([
    FullyConnectedLayer(n_in=training_x.get_value().shape[1], n_out=n_neurons),\
    FullyConnectedLayer(n_in=n_neurons, n_out=training_y.get_value().shape[1]),\
        mini_batch_size)
net.SGD(training_x, training_y, 1000, eta_choice)

# increase counter
loop_count += 1

#-----
# after a good initial point is found, we proceed to the extensive training

# initiate the deviation truncation criterion
med_deviante = 1000.

# loop until the deviation is smaller than the chosen truncation criterion
# we also truncate if the step size has become too small
while (med_deviante > trunc_diff) and (loop_count < max_iter) and (eta_choice/step_devide > min_e

    # continue to train the network if it has not converged yet
    net.SGD(training_x, training_y, num_epochs_choice, eta_choice/step_devide)

    # increase counter to avoid infinite loop
    loop_count += 1

    # check if the current step size is too large, i.e. cost does not change much
    if np.abs(np.mean(net.cost_train[:100]) \
        - np.mean(net.cost_train[-100:]))/(np.mean(net.cost_train[:100])) < 0.01:

        # if so, we make the step size smaller
        step_devide = step_devide*2.

#-----
# this is the validation step

```

```

""" this is the validation step
# calculate the deviation between the analytic approximation vs. the training models
# in principle, we should consider validation models here
w_array_0 = net.layers[0].w.get_value().T
b_array_0 = net.layers[0].b.get_value()
w_array_1 = net.layers[1].w.get_value()[0]
b_array_1 = net.layers[1].b.get_value()[0]

predict_flux = act_func(np.sum(w_array_1*act_func(np.dot(w_array_0,labels).T\
+ b_array_0), axis=1) + b_array_1)

# remember to scale back the fluxes to the normal metric
### here we choose the maximum absolute deviation to be the truncation criterion ###
med_deviate = np.max(np.abs((predict_flux - spectra[pixel_no,:])/0.8))

# return the trained network for this pixel
return net

#=====
# train in parallel for all wavelength pixels
import time
start_time = time.time()
pool = Pool(num_CPU)
net_array = pool.map(train_pixel,range(spectra.shape[0]))
print time.time() - start_time

#-----
# extract neural network parameters as numpy array
# the first layer
w_array_0 = np.array([net_array[i].layers[0].w.get_value().T\
                      for i in range(spectra.shape[0])])
b_array_0 = np.array([net_array[i].layers[0].b.get_value()\
                      for i in range(spectra.shape[0])])

# the second layer
w_array_1 = np.array([net_array[i].layers[1].w.get_value()[0]\
                      for i in range(spectra.shape[0])])

```



```

b_array_1 = np.array([net_array[i].layers[1].b.get_value()[0]\
                      for i in range(spectra.shape[0])])

#-----
# save the neural network parameters and record how we have scaled the labels
np.savez("neural_network.npz",\
        w_array_0=w_array_0, w_array_1=w_array_1,\
        b_array_0=b_array_0, b_array_1=b_array_1,\
        x_max=x_max, x_min=x_min)

```

Check how well we can predict the fluxes of spectra.

Ab-initio calculated spectrum vs. neural network reconstruction of a spectrum.

```

In [1]: # import python packages
%matplotlib inline
import matplotlib.pyplot as plt
from matplotlib import rcParams
import numpy as np

#-----
# define plot properties
def rgb(r,g,b):
    return (float(r)/256.,float(g)/256.,float(b)/256.)

cb2 = [rgb(31,120,180), rgb(255,127,0), rgb(51,160,44), rgb(227,26,28), \
        rgb(166,206,227), rgb(253,191,111), rgb(178,223,138), rgb(251,154,153)]
rcParams['lines.linewidth'] = 1
rcParams['axes.color_cycle'] = cb2
rcParams['font.family'] = 'Bitstream Vera Sans'
rcParams['font.size'] = 35

```

```
#=====
# restore trained neural network
temp = np.load("neural_network.npz")

w_array_0 = temp["w_array_0"]
w_array_1 = temp["w_array_1"]
b_array_0 = temp["b_array_0"]
b_array_1 = temp["b_array_1"]
x_min = temp["x_min"]
x_max = temp["x_max"]

# activation function
# make sure this function is consistent with the trained network
def act_func(z):
    return 1.0/(1.0+np.exp(-z))

#=====
# restore all spectra
temp = np.load("kurucz_spectra.npz")
wavelength = temp["wavelength"]
spectra = temp["Y_u_all"].T
labels = temp["labels_array"]

#-----
### testing spectra (not used in the training steps) ###
spectra_testing = spectra[:,1::2]
labels_testing = labels[:,1::2]

# scale labels as before
labels_testing = ((labels_testing.T - x_min)*0.8/(x_max-x_min) + 0.1).T

# predict flux using the neural network
predict_flux_array = []
```

```

# for each of the spectra
for i in range(labels_testing.shape[1]):

    # we predict the flux ...
    predict_flux_array.append(act_func(np.sum(w_array_1*(act_func(np.dot(w_array_0,\
        labels_testing[:,i]) + b_array_0)), axis=1) + b_array_1))

# convert into a numpy array
predict_testing = np.array(predict_flux_array).T

# recall that the predicted fluxes are scaled to [0.1,0.9]
# so here we have to rescale it back to the normal metric
predict_testing = (predict_testing - 0.1)/0.8

#-----
### we perform the same for training spectra ###
spectra_training = spectra[:,0::2]
labels_training = labels[:,0::2]
labels_training = ((labels_training.T - x_min)*0.8/(x_max-x_min) + 0.1).T
predict_flux_array = []
for i in range(labels_training.shape[1]):
    predict_flux_array.append(act_func(np.sum(w_array_1*(act_func(np.dot(w_array_0,\
        labels_training[:,i]) + b_array_0)), axis=1) + b_array_1))
predict_training = np.array(predict_flux_array).T
predict_training = (predict_training - 0.1)/0.8

#=====
# initiate the plot
fig = plt.figure(figsize=[22,50]);
ax = fig.add_subplot(111)
ax.spines['top'].set_color('none')
ax.spines['bottom'].set_color('none')
ax.spines['left'].set_color('none')
ax.spines['right'].set_color('none')
ax.tick_params(labelcolor='w', top='off', bottom='off', left='off', right='off')

```

“ . . . ”

```

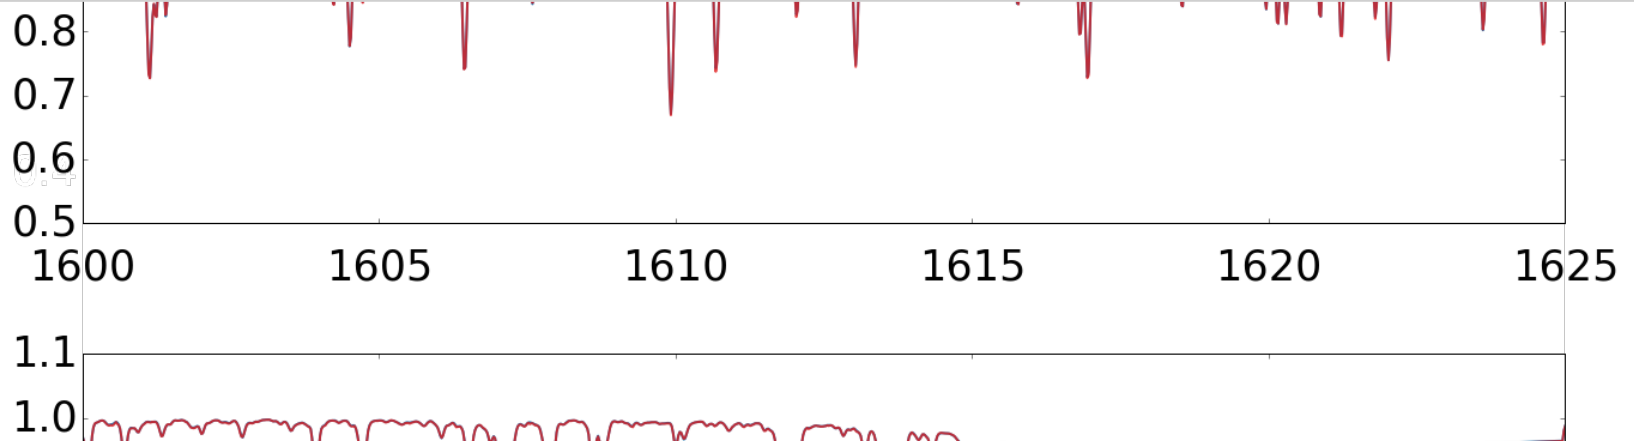
# axis labels
ax.set_xlabel("Wavelength [nm]", labelpad=30, fontsize=50);
ax.set_ylabel("Normalized Flux", labelpad=30, fontsize=50);

#-----
# break wavelength into different subplots
for i in range(8):
    ax = fig.add_subplot(8,1,i+1)
    plt.xlim([1500+i*25.,1500+(i+1)*25])
    plt.ylim([0.5,1.1])
    ax.tick_params(axis='x', pad=20);
    plt.plot(wavelength/10, spectra_testing[:,0],\
             color=cb2[0], lw=2, label="Ab-Initio Calculation")
    plt.plot(wavelength/10, predict_testing[:,0],\
             color=cb2[3], lw=2, label="The Payne Reconstruction", alpha=0.8)

#-----
# plot legend
plt.legend(loc="lower right", fontsize=35, frameon=False,\
          borderpad=0.2, labelspacing=0.2, scatterpoints=1)

# save figure
plt.tight_layout()
plt.savefig("flux_prediction.png")

```





Evaluate the median absolute deviation of all wavelength pixels.

```
In [2]: # initiate plot
fig = plt.figure(figsize=[18,12]);
ax = fig.gca();

# axis labels
ax.set_xlabel(r"Median Absolute Deviation");
ax.set_ylabel(r"Histogram of Pixels");

# plotting range
plt.xlim([0,0.003])

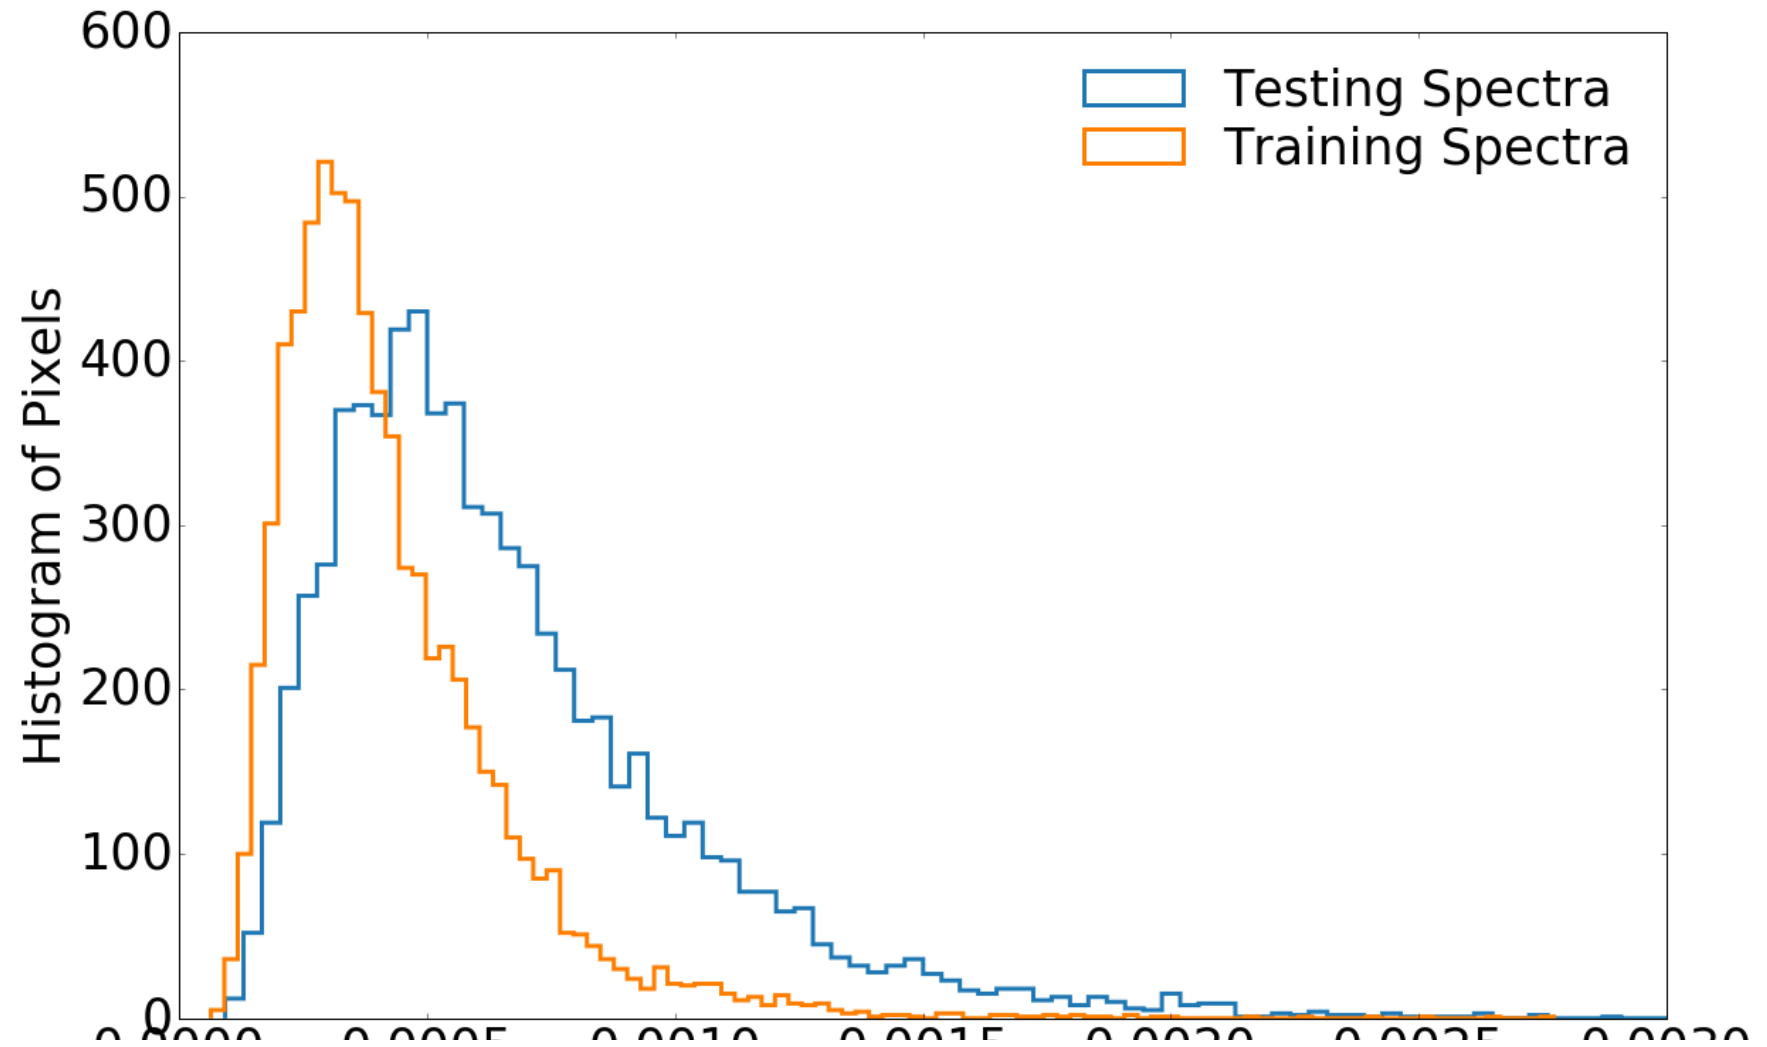
#-----
# plotting histogram
plt.hist(np.median(np.abs(spectra_testing-predict_testing), axis=1), bins=100,\
         histtype="step", lw=3, label="Testing Spectra", )
plt.hist(np.median(np.abs(spectra_training-predict_training), axis=1), bins=100,\
         histtype="step", lw=3, label="Training Spectra")

print np.median(np.abs(spectra_testing-predict_testing), axis=1).shape

#-----
# plot legend
plt.legend(loc="upper right", fontsize=35, frameon=False,\
```

```
borderpad=0.2, labelspace=0.2, scatterpoints=1)  
  
# save figure  
plt.tight_layout()  
plt.savefig("flux_deviation.png")
```

(7214,)



0.0000 0.0005 0.0010 0.0015 0.0020 0.0025 0.0030

Median Absolute Deviation

Fitting spectra with the trained network.

```
In [5]: # import python packages
import numpy as np
from multiprocessing import Pool
from scipy.optimize import curve_fit

#####
# number of processors for parallel computing
num_CPU = 4

#-----
# restore testing spectra
temp = np.load("kurucz_spectra.npz")
spectra = temp["Y_u_all"].T
labels = temp["labels_array"]
spectra_testing = spectra[:,1::2]
labels_testing = labels[:,1::2]

# scale the flux to the neural network scale
# if the fluxes have uncertainties, also remember to scale the uncertainties
spectra_testing = spectra_testing*0.8 + 0.1

#-----
# restore trained neural network
temp = np.load("neural_network.npz")
w_array_0 = temp["w_array_0"]
w_array_1 = temp["w_array_1"]
b_array_0 = temp["b_array_0"]
```

```

w_array_0 = temp[ "w_array_0" ]
b_array_1 = temp[ "b_array_1" ]
x_min = temp[ "x_min" ]
x_max = temp[ "x_max" ]

#####
# activation function
# make sure this function is consistent with the trained network
def act_func(z):
    return 1.0/(1.0+np.exp(-z))

# predict spectrum
def predict_spec(input_param, *labels):
    predict_flux = act_func(np.sum(w_array_1*(act_func(np.dot(w_array_0,labels)\
                                                                + b_array_0)), axis=1) + b_array_1)

    return predict_flux

# define function to perform testing steps in batch (over many spectra)
def fit_spectrum(spec_no):
    try:
        # here we set the noise (sigma) to be 0.01 dex
        # but note that we assume noiseless spectra in our testing case
        # here we also start with the right initial point
        # in practice multiple resets might be needed
        popt, pcov = curve_fit(predict_spec, [spec_no], spectra_testing[:,spec_no],\
                                p0 = (labels_testing[:,spec_no]-x_min)*0.8/(x_max-x_min) + 0.1,\
                                sigma=0.01,\
                                absolute_sigma=True, bounds=(0,1))

    except:
        # if the minimization does not converge, return some nul values
        popt = np.zeros(labels_testing.shape[0]) - 9999.
    return popt

#-----
# fit spectra in parallel
pool = Pool(num_CPU)
best_fit_results = np.array(pool.map(fit_spectrum.range(spectra_testing.shape[1]))).T

```



```

# recall that we have scaled the labels to train the network
# here we rescale the best-fitting results back to the original scale
best_fit_results = ((best_fit_results.T-0.1)*(x_max-x_min)/0.8+x_min).T

```

Plot the best fitting results vs. input parameters

```

In [4]: # import python packages
%matplotlib inline
import matplotlib.pyplot as plt
from matplotlib import rcParams
import numpy as np

#-----
# define plot properties
def rgb(r,g,b):
    return (float(r)/256.,float(g)/256.,float(b)/256.)

cb2 = [rgb(31,120,180), rgb(255,127,0), rgb(51,160,44), rgb(227,26,28), \
        rgb(166,206,227), rgb(253,191,111), rgb(178,223,138), rgb(251,154,153)]
rcParams['lines.linewidth'] = 1
rcParams['axes.color_cycle'] = cb2
rcParams['font.family'] = 'Bitstream Vera Sans'
rcParams['font.size'] = 35

#=====
# initiate the plot
fig = plt.figure(figsize=[25,9]);
ax = fig.add_subplot(111)
ax.spines['top'].set_color('none')
ax.spines['bottom'].set_color('none')

```

```

ax.spines['left'].set_color('none')
ax.spines['right'].set_color('none')
ax.tick_params(labelcolor='w', top='off', bottom='off', left='off', right='off')

# axis labels
ax.set_xlabel("Input Parameter", fontsize=40, labelpad=30);
ax.set_ylabel("Best Estimate", fontsize=40, labelpad=50);

#-----
# label name
label_name = [ "$\mathregular{T_{\mathrm{eff}}};[K]$", \
               "$\mathregular{\log\,g}$", \
               "[Fe/H]" ]

#-----
# loop over all labels
for u1 in range(len(label_name)):
    ax = fig.add_subplot(1, 3, u1+1)
    ax.tick_params(axis='x', pad=10);
    plt.locator_params(nbins=5)

    # plot results
    plt.scatter(labels_testing[u1,:], best_fit_results[u1,:], color="black", s=100)

#-----
# plotting parameters
# for Teff
if u1 == 0:
    # plotting range
    plt.xlim([2000,8000])
    plt.ylim([2000,8000])

    # plot 1-to-1 line guide line
    plt.plot([2000,8000],[2000,8000], lw=5, color=cb2[3])

    # plot the name of the label
    plt.text(2200, 7000, label_name[u1], fontsize=50)

```

```

# calculate and write the 1 sigma error
plt.text(4300, 2300, r"$\mathregular{\sigma}$ = "\
        + "%.1f" % (np.percentile(labels_testing[u1,:]-best_fit_results[u1:], 68.3)\
        - np.percentile(labels_testing[u1,:]-best_fit_results[u1:], 31.7)) + "K", fontsize=

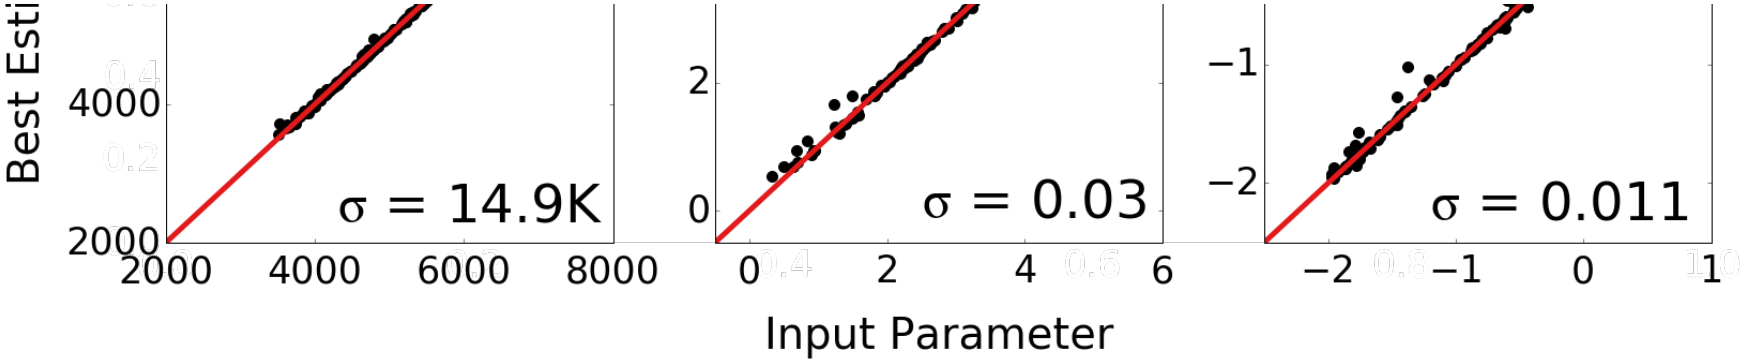
# forlogg
if u1 == 1:
    plt.xlim([-0.5,6.0])
    plt.ylim([-0.5,6.0])
    plt.plot([-0.5,6.0],[-0.5,6.0], lw=5, color=cb2[3])
    plt.text(-0.1, 4.9, label_name[u1], fontsize=50)
    plt.text(2.5, -0.1, r"$\mathregular{\sigma}$ = "\
            + "%.2f" % (np.percentile(labels_testing[u1,:]-best_fit_results[u1:], 68.3)\
            - np.percentile(labels_testing[u1,:]-best_fit_results[u1:], 31.7)), fontsize=50)

# for [Fe/H]
if u1 == 2:
    plt.xlim([-2.5,1.0])
    plt.ylim([-2.5,1.0])
    plt.plot([-2.5,1.0],[-2.5,1.0], lw=5, color=cb2[3])
    plt.text(-2.3, 0.4, label_name[u1], fontsize=50)
    plt.text(-1.2, -2.3, r"$\mathregular{\sigma}$ = "\
            + "%.3f" % (np.percentile(labels_testing[u1,:]-best_fit_results[u1:], 68.3)\
            - np.percentile(labels_testing[u1,:]-best_fit_results[u1:], 31.7)), fontsize=50)

#-----
# save figure
plt.tight_layout(w_pad=0.7,h_pad=0.7)
plt.savefig("fitting_results.png")

```





In []: